



First-order EMA approximation: per-token error metrics over training

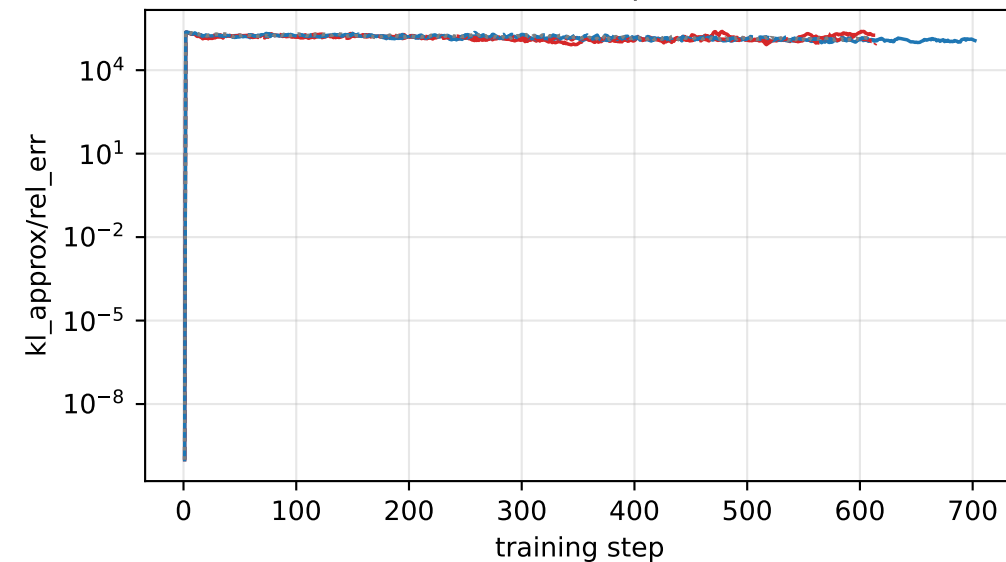
MAE per token



signed bias per token



relative error (eps=1e-8)



relative error (clamp [-10,10])

