

PHASE-A RESULTS: THE FREE EMA-KL SURROGATE PREVENTS STALENESS COLLAPSE, AND ITS THEORY PICKS THE RIGHT DOSE

BRIAN BARTOLDSON (EXPERIMENTS RUN AND REPORT DRAFTED WITH CLAUDE)

ABSTRACT. Asynchronous RL trains on rollouts from stale policy snapshots; the free EMA-KL surrogate $c(\log T_n - \log T_{n-\Delta})$ with $c = \alpha/(\Delta(1 - \alpha))$ regularizes against an implicit EMA anchor at zero extra cost. On Qwen3-4B/Countdown (TBA, no centering) we find: (1) **without** the penalty, training *collapses* at high staleness — homogeneous $\Delta=32$ falls from eval 0.75 to 0.14 by step 200, and a heterogeneous mixture $\Delta_i \in \{1, 4, 10, 32\}$ collapses ~ 60 steps later ($0.75 \rightarrow 0.002$ by step 300); (2) **with** the penalty at the theory-prescribed strength, every arm is stable and reaches eval 0.79–0.83; (3) per-rollout calibration $c_i = \alpha/(\Delta_i(1 - \alpha))$ matches the *best* uniform dose found by a $3\times$ sweep (peaks 0.816 vs. 0.817) — calibration’s value is that the theory selects the right strength *a priori*, without a sweep; the $3\times$ larger uniform dose costs ~ 1.5 points. A homogeneous $\Delta=10$ reference through the new machinery reproduces the historical baseline (0.829 vs. 0.827–0.831), and per- Δ telemetry confirms the surrogate’s predicted small- Δ variance (rel. error 5.9 at $\Delta \leq 2$ vs. 0.19 at $\Delta \geq 21$) without harming training. NVFP4-emulated symmetric training is stable at smoke scale.

1. SETUP

All runs: Qwen3-4B-Instruct-2507 on Countdown with TBA (no centering, importance sampling on), lr 10^{-6} , $K=16$, seq. 4096, 720 target steps, eval every 20 steps, $\alpha = 0.9$, 2 nodes ($4\times H100$ inference / $4\times H100$ training); 24 h allocations (curves end at steps 490–690). Configs in Appendix A.

Heterogeneous staleness. `--rl.staleness-offsets` $o_{0..3}$ pins each of 4 data-parallel vLLM ranks to a fixed checkpoint lag (reload every step to step $-o_r$). K -groups stay policy-homogeneous; training batches pool ranks and therefore mix policy ages (verified in the rollout parquets: the trainer’s step-10 batch carries generating checkpoints $\{9, 6, 0, 0\}$). The homogeneous arms use offsets 10, 10, 10, 10 (reference) or the legacy reload rule at `async_level=32` ($\beta=0$ control).

Anchors and calibration. Every arm anchors both the importance ratio and the KL term on the sampler’s own vLLM logprobs. Per-rollout calibration uses each sample’s true lag: $c_i = \alpha/(\Delta_i(1 - \alpha))$, $\Delta_i \geq 1$; the uniform arms apply one dose $w = \beta c_{\text{glob}}$ to every rollout ($c_{\text{glob}} = 0.9/3.2 = 0.28125$). The per-token clamp is in raw log-ratio units (± 10 before the coefficient), so truncation is Δ -invariant.

2. RESULTS

Finding 1: staleness alone collapses training; the free surrogate prevents it. The homogeneous $\Delta=32$, $\beta=0$ control peaks at eval 0.752 (step 100), then falls off a cliff between steps 160–180 (train reward $0.68 \rightarrow 0.19$; eval 0.14 at step 200) and oscillates around 0.2–0.45 for 400 more steps without recovering. The heterogeneous $\beta=0$ control — average lag $4\times$ smaller — shows the same failure delayed by ~ 60 steps, preceded by a characteristic train-reward overshoot (0.94 at step 127), then eval 0.002 by step 300. Every arm with the EMA-KL surrogate on, at any tested strength, trains stably through and far beyond these windows. This is BF16 end to end: staleness alone, no quantization, breaks unregularized async RL at these lags, and the zero-cost surrogate fixes it.




FIGURE 1. Collapse without the KL penalty; stability with it. Yellow (heterogeneous) and violet (homogeneous $\Delta=32$) are $\beta=0$: both collapse (steps ~ 230 and ~ 170) and never recover. Aqua (uniform dose $w=0.0045$) trains through the same window untouched; blue is the homogeneous $\Delta=10$ reference.




FIGURE 2. All stable arms. Calibrated per-rollout c_i (aqua) vs. the two uniform doses (yellow $w=0.0045$, violet $w=0.015$) under the same $\Delta_i \in \{1, 4, 10, 32\}$ mixture; blue is the homogeneous reference.

Finding 2: calibration matches the best swept dose — without the sweep. Peaks: calibrated 0.816 (step 520), uniform $w=0.0045$ 0.817 (step 500), uniform $w=0.015$ 0.801; reference 0.829. The calibrated arm and the best uniform arm are statistically indistinguishable at one seed, and both sit ~ 1 point under the homogeneous reference. Two honest readings. First, on this task and Δ spread, per-rollout weighting adds no *performance* over the best *tuned* uniform dose — the group-summed penalty appears to care mostly about total strength. Second, the best uniform dose is exactly the one the theory prescribes ($w = \beta_{\text{ref}} c_{\text{ref}} = 0.0045$, transported from the known-good homogeneous run via the βc invariance), and tripling it costs 1.5 points: the

calibration theory’s practical payoff here is *dose selection without a sweep*, which is also what a practitioner changing Δ (or mixing Δ_i) actually needs.




FIGURE 3. Relative error of the surrogate vs. exact EMA-KL by lag stratum (calibrated run; window-median smoothed; log scale).

Finding 3: the surrogate’s error is strongly Δ -dependent — and training doesn’t care. Median rel. error of means: 5.88 ($\Delta \leq 2$), 2.13 ($[3, 6]$), 0.48 ($[7, 20]$), 0.19 (≥ 21) — the theory note’s finite-difference variance warning, measured. Despite a quarter of tokens carrying rel.-error-~6 estimates, the calibrated arm matched the reference; the summed, β -scaled penalty is robust to zero-mean estimator noise. Suggested refinement (untested): floor Δ_i at ~ 4 in c_i .

Finding 4: mechanism and baseline controls. The homogeneous $\Delta=10$ reference through the full new machinery (offsets mechanism, vLLM anchor, raw-units clamp) peaks 0.829 / final 0.828 — inside the historical no-centering band (0.827–0.831). NVFP4-emulated symmetric training (QDQ on trainer forwards and sampler checkpoints, bit-exact) is stable at 60-step smoke scale; checkpoints verified on-grid. FP4 collapse/rescue arms are the next queue batch.

3. NEXT STEPS

Paired seeds (≥ 3) on calibrated vs. best-uniform; FP4 $\beta=0$ ladder and rescue arms (does quantization shift the collapse boundary or the required dose?); Δ_i floor test; a second task (MATH). Deferred items and operational cautions are recorded in `EXPERIMENT_ROADMAP.md` §7 and `report_fp4/OPERATIONS_LOG.md`.

APPENDIX A. CONFIGURATIONS

Common: model=Qwen/Qwen3-4B-Instruct-2507, LR=1e-6, sampled_k=train_k=16, steps=720, type=tb, conf=COUNTDOWN.toml, seq_length=4096, eval_steps=20, reference_mode=ema, ema_alpha=0.9, kl_centering=false, vllm_logprobs=true, batch 512, wandb fp4_KL, branch fp4_KL.

Arm	β	async	offsets	Δ -source
homog10_approxKL_b005_vllmLP	0.005	10	10,10,10,10	per_rollout
hetero_approxKL_b005_perRollout	0.005	32	1,4,10,32	per_rollout
hetero_approxKL_global_w0045	0.016	32	1,4,10,32	global
hetero_approxKL_global_w015	0.0533	32	1,4,10,32	global
hetero.b0_async32	0	32	1,4,10,32	—
bf16.b0_async32	0	32	— (legacy rule)	—
fp4sym.smoke (60 steps)	0.005	10	—	global

KL penalty: $-\beta \sum_t c \text{clamp}(\log T_n - \log T_{n-\Delta}, \pm 10)$ added to advantages (no centering); c per-token (c_i) or global per the table. FP4 arms add `rollout_quant=nvfp4` + `fake_quant=nvfp4` (252 patched linears; embeddings, `lm_head`, norms BF16).

APPENDIX B. EVAL TABLE (SELECTED STEPS)

Step	100	200	300	400	500	final
Homog. $\Delta=10$ reference	0.731	0.754	0.773	0.799	0.806	0.828 (680)
Hetero, per-rollout c_i	0.725	0.764	0.811	0.799	0.815	0.810 (540)
Hetero, uniform $w=0.0045$	0.722	0.759	0.782	0.794	0.817	0.806 (560)
Hetero, uniform $w=0.015$	0.670	0.730	0.763	0.774	0.801	0.788 (560)
Hetero, $\beta=0$	0.725	0.662	0.002	0.123	—	0.222 (480)
Homog. $\Delta=32$, $\beta=0$	0.752	0.140	0.246	0.271	0.319	0.340 (540)